

Overview of Data Analytics: Data Mining & Warehousing

Vikram Pudi
vikram@iiit.ac.in
IIIT Hyderabad

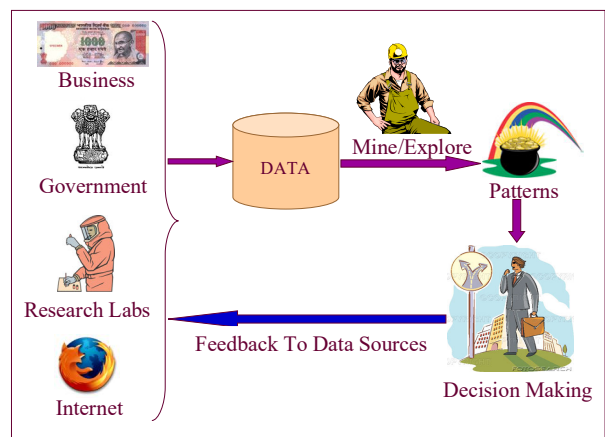
Originated from DB community...

- Traditional Database Systems
 - Indexing
 - Query languages
 - Query optimization
 - Transaction processing
 - Recovery ...
- XML, Semantic web
- OO and OR DBMS ...
- *Data Mining*

2

Data Mining

Automated extraction of
interesting patterns from large
databases



4

Types of Patterns

- **Associations**
 - *Coffee* buyers usually also purchase *sugar*
- **Clustering**
 - Segments of customers requiring different promotion strategies
- **Classification**
 - Customers expected to be *loyal*



Association Rules

That which is infrequent is not
worth worrying about.

6

Association Rules

Transaction ID	Items
1	Tomato, Potato, Onions
2	Tomato, Potato, Brinjal, Pumpkin
3	Tomato, Potato, Onions, Chilly
4	Lemon, Tamarind

Rule: Tomato, Potato \rightarrow Onion (confidence: 66%, support: 50%)

Support(X) = |transactions containing X| / |D|

Confidence(R) = support(R) / support(LHS(R))

Problem proposed in [AIS 93]: Find all rules satisfying user given minimum support and minimum confidence.

7


Association Rule Applications

- E-commerce
 - People who have bought *Sundara Kandan* have also bought *Srimad Bhagavatham*
- Census analysis
 - Immigrants are usually male
- Sports
 - A chess end-game configuration with "white pawn on A7" and "white knight dominating black rook" typically results in a "win for white".
- Medical diagnosis
 - Allergy to *latex rubber* usually co-occurs with allergies to *banana* and *tomato*

8

The Apriori Algorithm

Idea: An itemset can be frequent only if all its subsets are frequent.



Apriori(DB, minsup):
 $C = \{\text{all 1-itemsets}\}$
 // candidates = singletons
 while (|C| > 0):
 make pass over DB, find counts of C
 $F = \{\text{sets in } C \text{ with count} \geq \text{minsup} \cdot |DB|\}$
 output F
 $C = \text{AprioriGen}(F)$ // gen. candidates

AprioriGen(F):
 for each pair of itemsets X, Y in F:
 if X and Y share all items, except last
 $Z = X \cup Y$ // generate candidate
 if any imm. subset of Z is not in F:
 prune Z // Z can't be frequent

9

Types of Association Rules

- Boolean association rules
- Hierarchical rules
 - stationary
 - pens
 - reynolds
 - cross
 - pencils
 - natraj
 - steadler
 - reynolds \rightarrow pencils
- Quantitative & Categorical rules
 - (Age: 30...39), (Married: Yes) \rightarrow (NumCars: 2)

10

More Types of Association Rules

- Cyclic / Periodic rules
 - Sunday \rightarrow vegetables
 - Christmas \rightarrow gift items
 - Summer, rich, jobless \rightarrow ticket to Hawaii
- Constrained rules
 - Show itemsets whose average price > Rs.10,000
 - Show itemsets that have television on RHS
- Sequential rules
 - Star wars, Empire Strikes Back \rightarrow Return of the Jedi

11

Classification



To be or not to be: That is the question.

- William Shakespeare

12

The Classification Problem

Outlook	Temp (°F)	Humidity (%)	Windy?	Class
sunny	75	70	true	play
sunny	80	90	true	don't play
sunny	85	85	false	don't play
sunny	72	95	false	don't play
sunny	69	70	false	play
overcast	72	90	true	play
overcast	83	78	false	play
overcast	64	65	true	play
overcast	81	75	false	play
rain	71	80	true	don't play
rain	65	70	true	don't play
rain	75	80	false	play
rain	68	80	false	play
rain	70	96	false	play
sunny	77	69	true	?
rain	73	76	false	?

Play Outside?

Model relationship between class labels and attributes

e.g. outlook = overcast \Rightarrow class = play

\Rightarrow Assign class labels to new data with *unknown* labels

13


Applications

- Text classification
 - Classify emails into spam / non-spam
 - Classify web-pages into yahoo-type hierarchy
 - NLP Problems
 - Tagging: Classify words into verbs, nouns, etc.
- Risk management, Fraud detection, Computer intrusion detection
 - Given the properties of a transaction (items purchased, amount, location, customer profile, etc.)
 - Determine if it is a fraud
- Machine learning / pattern recognition applications
 - Vision
 - Speech recognition
 - etc.
- All of science & knowledge is about predicting future in terms of past
 - So classification is a very fundamental problem with ultra-wide scope of applications

14

k-Nearest Neighbours


- Model = Training data
- Classify record R using the k nearest neighbours of R in the training data.
- Most frequent class among k NNs
- Distance function could be euclidean
- Use an index structure (e.g. R* tree) to find the k NNs efficiently




15

Decision Trees

Outlook	Temp (°F)	Humidity (%)	Windy?	Class
sunny	75	70	true	play
sunny	80	90	true	don't play
sunny	85	85	false	don't play
sunny	72	95	false	don't play
sunny	69	70	false	play
overcast	72	90	true	play
overcast	83	78	false	play
overcast	64	65	true	play
overcast	81	75	false	play
rain	71	80	true	don't play
rain	65	70	true	don't play
rain	75	80	false	play
rain	68	80	false	play
rain	70	96	false	play



16



Clustering

Birds of a feather flock together.

The Clustering Problem

Find groups of similar records.

Need a function to compute similarity, given 2 input records

Outlook	Temp (°F)	Humidity (%)	Windy?
sunny	75	70	true
sunny	80	90	true
sunny	85	85	false
sunny	72	95	false
sunny	69	70	false
overcast	72	90	true
overcast	73	88	true
overcast	64	65	true
overcast	81	75	false
rain	71	80	true
rain	65	70	true
rain	75	80	false
rain	68	80	false
rain	70	96	false

\Rightarrow Unsupervised learning

17

18

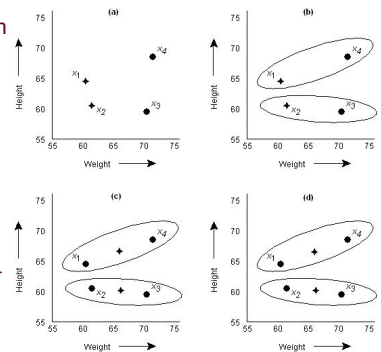
Applications

- Targetting similar people or objects
 - Student tutorial groups
 - Hobby groups
 - Health support groups
 - Customer groups for marketing
 - Organizing e-mail
- Spatial clustering
 - Exam centres
 - Locations for a business chain
 - Planning a political strategy

19

Partitioning technique: k-Means


- Initial k means = random records
- Iterate as long as clusters change:
 - Put each record X in the cluster to whose mean it is closest
 - Recompute means as the average of all points in each cluster



20

Data Warehousing

- Extract, transform, load data from multiple sources in an enterprise
- Provide unified view for top management
- OLAP server provides multi-dimensional view for manual exploration of patterns



21

Examples of OLAP

Comparisons (this period v.s. last period)

Show me the sales per store for this year and compare it to that of the previous year to identify discrepancies

Ranking and statistical profiles (top N/bottom N)

Show me sales, profit and average call volume per day for my 10 most profitable salespeople

Custom consolidation (market segments, ad hoc groups)

Show me an abbreviated income statement by quarter for the last four quarters for my northeast region operations

Take Home

- Data mining is a mature field
- Don't waste time developing new algorithms for core tasks
- Focus on applications to challenging kinds of data
 - Streams, Distributed data, Multimedia, Web, ...
- Most effort is in how to map domain problems to data mining problems
- And how to make sense of the output.

23



24