

Efficient Discovery of Concise Association Rules from Large Databases

Vikram Pudi
vikram@iiit.ac.in
IIIT Hyderabad

Talk Outline

- Introduction
- Mining Association Rules
- Conciseness of Mining Results
- Conclusions

Talk Outline

- Introduction
 - Define Association Rules
 - Applications
 - Types of Association Rules
 - Interestingness Measures
 - Privacy
- Mining Association Rules
- Conciseness of Mining Results
- Conclusions

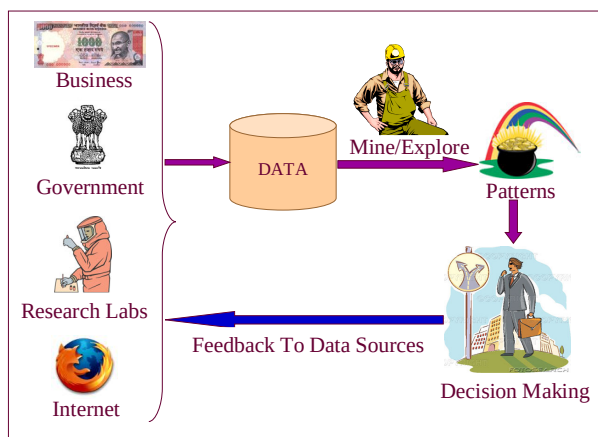
3

What is Data Mining?

Automated extraction of interesting patterns from large databases

Types of Patterns

- **Associations**
 - *Coffee* buyers usually also purchase *sugar*
- **Sequence Patterns**
 - After seeing *Superman*, people usually see *Star Wars*
- **Clustering**
 - Segments of customers requiring different promotion strategies
- **Classification**
 - Customers expected to be *loyal*



5

Association Rules

D :

Transaction ID	Items
1	Tomato, Potato, Onions
2	Tomato, Potato, Brinjal, Pumpkin
3	Tomato, Potato, Onions, Chilly
4	Lemon, Tamarind

Rule: Tomato, Potato \rightarrow Onion (confidence: 66%, support: 50%)

Support(X) = |transactions containing X| / |D|

Confidence(R) = support(R) / support(LHS(R))

Problem proposed in [AIS 93]: Find all rules satisfying user given minimum support and minimum confidence.


7

Typical Solution Strategy

- STEP 1: Find all *frequent* itemsets (computationally expensive)
 - Itemset X is frequent iff $\text{support}(X) \geq \text{minsup}$
- STEP 2: Find *rules* from the frequent itemsets (computationally inexpensive)
 - Rule quantity: too many rules are usually generated
 - Rule quality: not all rules are interesting

8

Difficulty



- Extremely computationally expensive
- Naïve solution
 - exponential time and memory w.r.t. |I|
 - linear time w.r.t. |D|
- Typically, |I| is in thousands, |D| is in billions...

9

Applications

- E-commerce
 - People who have bought *Sundara Kadam* have also bought *Srimad Bhagavatham*
- Census analysis
 - Immigrants are usually male
- Sports
 - A chess end-game configuration with "white pawn on A7" and "white knight dominating black rook" typically results in a "win for white".
- Medical diagnosis
 - Allergy to *latex rubber* usually co-occurs with allergies to *banana* and *tomato*

Killer Apps

Classification

Idea: Model of each class consists of frequent itemsets for that class. Compare new transactions with each model and select "nearest" one.

Advantages: Scales well to thousands of attributes and billions of rows.


Recommendation Systems

- People who listen to songs that you listen, have also listened to these other songs...
- People who have bought these books, have also bought these other books...

11

Types of Association Rules

- Boolean association rules
- Hierarchical rules



dhoti, saree \rightarrow t-shirt

- Quantitative & Categorical rules
 - (Age: 30...39), (Married: Yes) \rightarrow (NumCars: 2)

12

More Types of Association Rules

- Cyclic / Periodic rules
 - Sunday → vegetables
 - Christmas → gift items
 - Summer, rich, jobless → ticket to Hawaii
- Constrained rules
 - Show itemsets whose average price > Rs.10,000
 - Show itemsets that have television on RHS
- Sequential rules
 - Star wars, Empire Strikes Back → Return of the Jedi

13

Interestingness Measures

14

Traditional Measures

- Confidence: Likelihood of a rule being true
- Support:
 - Statistical significance: Data supports rule
 - Applicability: Rule with high support is applicable in large number of transactions

15

Problem with Confidence

- Researcher → Coffee (confidence: 90%)
- This rule intuitively means:
 - Researchers have a *strong tendency* to drink coffee.
- But if 95% of general population drinks coffee, then this is a **bad** rule.
- Solution: For, $X \rightarrow Y$,
 - Interest = $P(X, Y) / P(X) P(Y)$
 - Conviction = $P(X) P(\neg Y) / P(X, \neg Y)$
 - Reason: $X \rightarrow Y \Leftrightarrow \neg X \vee Y \Leftrightarrow \neg(X \wedge \neg Y)$

16

Surprising Patterns

- An itemset is uninteresting if its support can be **estimated** based on:
 - supports of its subsets
 - its support at earlier points in time
- Key Idea:
 - For an itemset X , remove items in each transaction that are not in X
 - If resulting database can be *compressed* well, then X is uninteresting (as it encodes less information)

17

Current Status of Interestingness

- *minsup* is required.
- So, get frequent itemsets first.
- Other interestingness measures can be applied later.
- Open Problem: How to select a good minimum support?

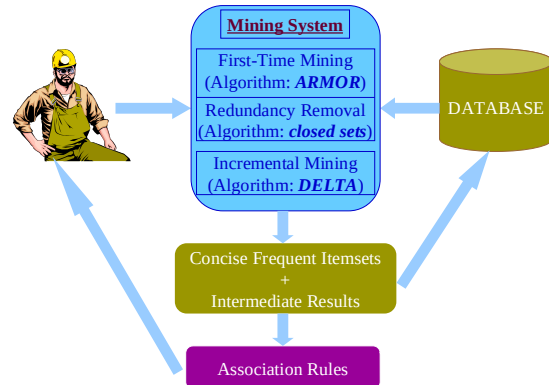
18

Issue: Privacy

- Users provide **inaccurate data** to protect their privacy.
- How can inaccurate data be effectively mined?
- How can data be **modified** in such a way as to ensure **data privacy** and rule accuracy?
- How can data be modified in such a way as to ensure **rule privacy**? – hide sensitive rules
- Can mined results be used to **retrieve original data** transactions?

19

Architecture for Association Rule mining

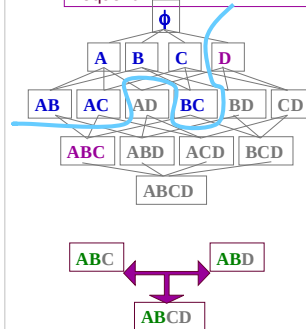


Talk Outline

- Introduction
- Mining Association Rules
 - Apriori
 - Partition
 - Sampling
 - Incremental Mining
 - FP-Growth (Frequent Pattern Growth)
 - Optimal Infeasible Algorithm: **Oracle**
 - ARMOR** (Association Rule Mining Based on **ORacle**)
- Conciseness of Mining Results
- Conclusions

The Apriori Algorithm

Idea: An itemset can be frequent only if all its subsets are frequent.




Apriori(DB, minsup):
 $C = \{\text{all 1-itemsets}\}$ // candidates = singletons
 while ($|C| > 0$):
 make pass over DB, find counts of C
 $F = \{\text{sets in } C \text{ with count} \geq \text{minsup} * |DB|\}$
 output F
 $C = \text{AprioriGen}(F)$ // gen. candidates

AprioriGen(F):
 for each pair of itemsets X, Y in F:
 if X and Y share all items, except last
 $Z = X \cup Y$ // generate candidate
 if any imm. subset of Z is not in F:
 prune Z // Z can't be frequent

22


Data Structures: For 1&2-itemsets



23

For Itemsets with length > 2

- Original Apriori uses **hash-trees**
 - Following **forest** structure (see figure) is as effective, and much simpler
 - Each leaf node has a **counter**
- Let current transaction $t = ACD$. Clearly, no child of B can be in t . So, the **search space is pruned**



Increment_count(t):
 for each item x in t # e.g. $x = A, C, D$
 $n = \text{root node cp. to } x$ # $n = A, C, D$
 $t_v = \text{bit-vector cp. to } t$ # e.g. 1011
 Increment_count(n, t_v)
Increment_count(n, t_v): # e.g. $n = A$
 # recursive depth-first search (DFS)
 if n is a leaf: $n.\text{count}++$
 for each child m of n # AB, AC, AD
 if item $m - n \in t_v$ # $m - n = B, C, D$
 Increment_count(m, t_v)

24

Analysis of Apriori

- ✓ Minimum number of candidates possible (more or less)
- ✗ **i/o intensive**: too many database scans
- ✗ **cpu intensive**: counting technique traverses data-structures for each transaction

25

The Partition Algorithm

Count of some itemset in each partition

$F_1 \ P_1 \ c_1 < minsup * |P_1|$
 $F_2 \ P_2 \ c_2 < minsup * |P_2|$
 $F_3 \ P_3 \ c_3 < minsup * |P_3|$
 $F_4 \ P_4 \ c_4 < minsup * |P_4|$
 $c < minsup * |DB|$
 Cannot be frequent overall.

- Conceptually partition horizontally
- Pass 1**: Mine each partition separately with same relative *minsup*. E.g. Use Apriori for this.
- Club individual mining results
- Pass 2**: Check if clubbed results are frequent in whole database
- An itemset infrequent in each partition will be infrequent.

26

Analysis of Partition

- ✓ Only two i/o scans over database
- ✗ Frequent itemset mining is mostly cpu-intensive
- ✗ Partition is **cpu intensive**
 - same counting technique as Apriori
 - Number of candidates in both scans similar to that of Apriori
 - So, does double the work
- ✗ Starts fresh Apriori for each partition instead of using previous partitions' results
- ✗ Sensitive to **skew** in data distribution
 - Some partitions may be very different and can contribute spurious candidates

27

Sampling

Pass 1: Count F_s and N_s over DB

F may contain itemsets not in F_s or N_s
How to find them?

- Promoted borders**: Itemsets in N_s that become frequent
- Observation 1**: If there are any promoted borders, then our estimate of F is wrong
- Observation 2**: Only supersets of the promoted borders can be frequent

Pass 2: Generate such supersets and find their counts over DB

- Do not generate itemsets which have infrequent subsets (i.e. $N_s - \{\text{promoted borders}\}$)

Idea: Mine a random sample S (of DB) to get F_s and N_s and treat F_s as an estimate of F

Negative Border N_s (of F_s):
Minimally infrequent itemsets
 \Rightarrow Infrequent candidates of Apriori
 \Rightarrow Has no infrequent subsets

28

Analysis of Sampling

- ✓ Estimate of frequent itemsets is usually quite accurate
 - Exact frequent itemsets not needed in many applications
- ✗ Sampling itself may require one database scan if random access is not possible (due to lack of index)
- ✗ If exact frequent itemsets are required, the cpu-cost is more than Apriori because
 - same counting technique as Apriori
 - first scan counts similar number of candidates as Apriori
 - cpu-cost of second scan is extra

29

Incremental Mining

Database = $DB \cup db$
 DB – original database
 db – increment

Find rules for $DB \cup db$ and for db

Input: F^{DB}, N^{DB}

Output: $F^{DB \cup db}, N^{DB \cup db}, F^{db}, N^{db}$

Idea 1: Treat DB and db as two partitions. Apply partition technique.

Idea 2: Treat DB as a sample. Apply sampling (negative border) technique.

Idea 3: Idea 1 + Idea 2

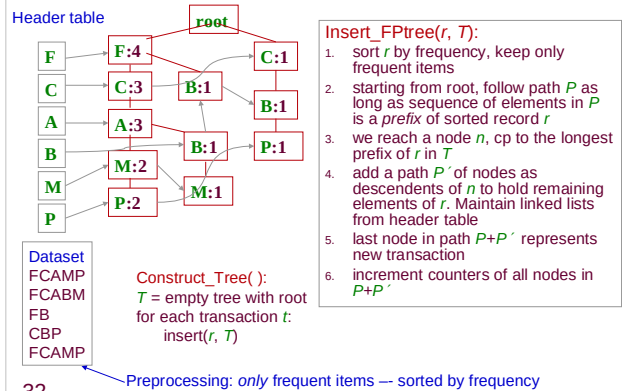
30

Practice Problem

Show the steps of Apriori on the following dataset with mincount = 3. Clearly show candidate and frequent itemsets of each length. Dataset: facdgmip, abcfimo, bfhjo, bcksp, afclpmn.

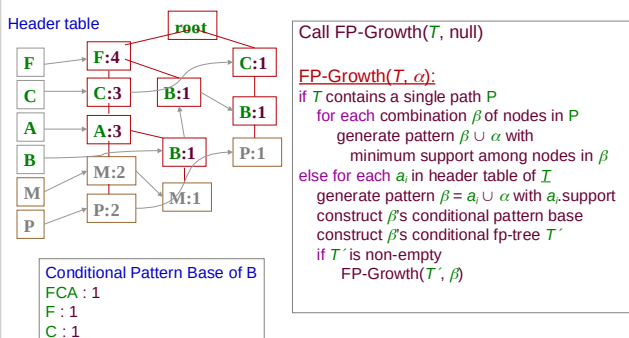
31

FP-Growth: Growing the FP-Tree



32

FP-Growth: Mining the FP-Tree



33


Analysis of FP-Growth

- ✓ Reuses work done for processing transactions that share a common prefix
- ✓ Counting technique is thus different from Apriori
 - ✓ No explicit data-structure traversal for each transaction
- ✓ Very effective for *dense* datasets
- ✗ Not effective for *sparse* datasets
 - ✗ If a sequence of items appears in even one transaction, it will be represented in the fp-tree
 - ✗ For a large random sparse dataset, every sequence is likely to appear in at least one transaction
 - ✗ The fp-tree then grows linearly with database size
- ✗ Invalid claims of "no candidates"
 - ✗ Every candidate that is counted in Apriori is also counted in FP-growth (and a counter is maintained for it)

34

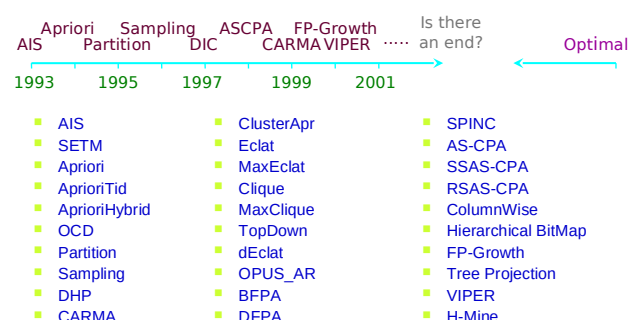
Problem Formulation

- Input
 - Database D
 - Minimum support threshold $minsup$
- Output:
 - Frequent Itemsets F : Itemsets whose support $\geq minsup$
 - Negative Border N : Minimally infrequent itemsets
 - Supports of $F \cup N$



35


Feeding Frenzy




Optimal Algorithm: Oracle

- Magically knows identities of frequent itemsets before mining begins. Therefore, has to only determine the counts of these itemsets in *one* pass over the database
- Minimum work required from any algorithm
- Careful design of data structures to ensure optimal access and enumeration of itemsets

Counting 1&2-itemsets



Counting Longer Itemsets ($k > 2$)




- Partition database
- For each itemset, compute the list of transaction-ids (*tidlist*) containing it
- Initiate *tidlist intersections* from frequent singletons
- Depth-first traversal
- Optimize using *tid-vector* approach

Tidset Intersection

Tid-vector V : 0010101101....

Tid-list L : 5, 6, 10, 29, ...



- Cost of intersection = $\Theta(|L|)$
- Cost of *tid-vector* construction
 - Proportional to number of "1"s in V
 - Amortized over many intersections
 - Space for V can be statically allocated

No wasted Enumeration

- All 1-itemsets are either frequent or in -ve border
- Only combinations of *frequent* 1-itemsets enumerated for *pairs*
- Depth-first search ensures each itemset is visited only *once*

Enumeration Cost = $\Theta(1)$

- Direct lookup arrays for 1&2-itemsets. Best in unit-cost RAM model
- For longer itemsets, cost = $\Theta(|X.childset|)$ resulting in $\Theta(1)$ cost per itemset overall
- All operations involve array and pointer lookups, which cannot be improved upon

Oracle Features

- Uses **direct lookup arrays** for 1-itemsets and 2-itemsets
- Uses **DAG** structure for longer itemsets
- No wasted enumeration of itemsets
- Enumeration cost per itemset = $\Theta(1)$
- **Caveat:** Not really optimal
 - Doesn't share work for transactions that are significantly similar. E.g. if 2 transactions are identical, it does the same work for both

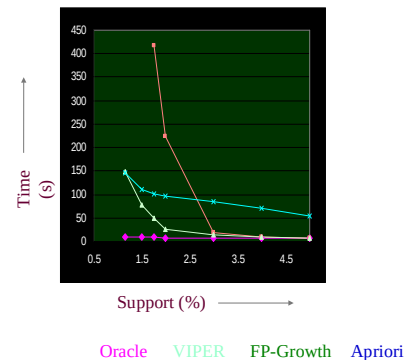
43

Performance of Current Algorithms

Performance Setup

- **Algorithms:** Oracle, VIPER, FP-growth, Apriori
- **Variety of Databases**
 - File-system backend
 - Integration with commercial RDBMS
 - Cache data to file-system and run algorithm
 - Implement algorithm as stored procedure
 - Implement algorithm in SQL
- Extreme and typical values of *minsup*

Response Times of Current Algorithms



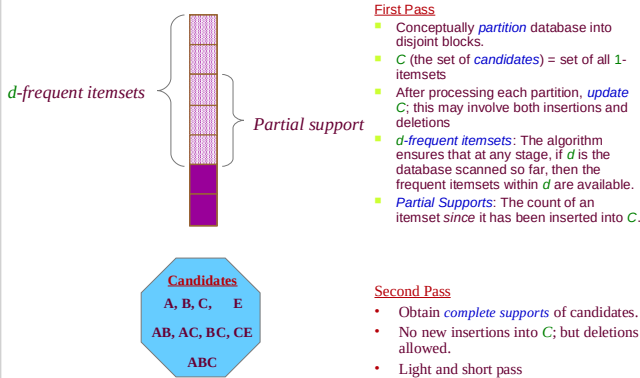
Online Algorithm

ARMOR: Association Rule
Mining based on ORacle

ARMOR


- *Minimal* changes to Oracle
- Maximum *two* passes over database
- “Short and light” second pass
- **Performance:** Within *twice* of Oracle for a variety of real and synthetic databases

ARMOR Processing



Candidate Generation

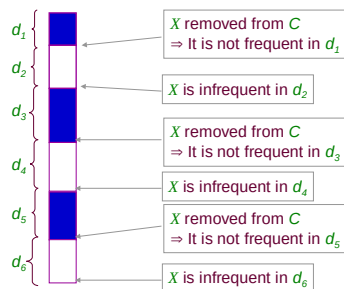
Itemsets can move freely between being partially-frequent, negative border and partially-infrequent.



Observation: An itemset can become partially frequent iff it has some subset in N which moves to F . Such itemsets are called *promoted borders*.

Proof of Correctness

Consider the life of an itemset X in the set of candidates, C
Solid area represents that X was in C
Blank area represents that X was not in C

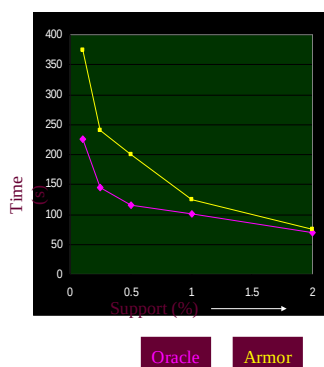


Since X is infrequent in every block, it is infrequent overall.


51

Performance of ARMOR

Response Times of ARMOR



Memory Utilization of ARMOR



54

Conclusions: Is ARMOR++ Feasible?

- Number of **candidates** in ARMOR are only **10%** more than minimum (all frequent itemsets+negative border)
- Number of **passes** is effectively less than **two**
- So, scope for improvement appears to be limited ...
- **Caveat:** Doesn't share work for transactions that are significantly similar. E.g. if 2 transactions are identical, it does the same work for both

Talk Outline

- Introduction
- First-time Mining of Association Rules
- Conciseness of Mining Results
 - Background
 - Closed Itemsets Framework
- Conclusions

Problem: Too many rules!

Dataset	<i>minsup</i>	#frequent
Sparse	0.1%	27,532
Dense	70%	48,969

Most are redundant

Post-mining Rule Pruning Schemes

- E.g. Output only rules that satisfy a user-given **minimum improvement**.
 - **Improvement:** Minimum difference between confidence of a rule and any of its sub-rules with the same RHS.
- **Problem:** What if mining itself is infeasible due to large output size?

58

Constrained Association Rules

- User specifies constraints on what kind of rules he is looking for. E.g. **RHS should contain milk**.
- **Problem:**
 - User may not have any constraints in mind.
 - Artificial, if purpose is just to reduce number of rules.

59

Maximal Frequent Itemsets

- A maximal frequent itemset is one that has no frequent supersets. (Also, called **positive border**.)
- **Problems:**
 - Identity of subsets can be deduced, but not their supports.
 - Subsets may have unexpected supports and thus be interesting on their own.
 - Cannot be used to form rules, since supports of subsets is necessary for this.

60


Closed Itemsets [Zak00]

Closed Itemsets Definition

- Tidset of an itemset X :
 - $t(X)$ = set of tids of transactions containing X
- Itemset of a tidset T :
 - $i(T)$ = items common to all transactions in T
- Closure Operator $c(X) = i(t(X))$:
 - Extension: $X \subseteq c(X)$
 - Monotonicity: If $X \subseteq Y$, then $c(X) \subseteq c(Y)$
 - Idempotency: $c(c(X)) = c(X)$
- Closed Itemset X : Iff $c(X) = X$

62

Closed Itemsets Redefinition



support(A) = support(AB)
support(AC) = support(ACD)
Etc.

An itemset is closed iff it has no superset with same support.

63

Mining Frequent Closed Itemsets

In any algorithm for frequent itemset mining:

- If an itemset has a subset with same support, don't generate any of its supersets as candidates.
- E.g. in Apriori, remove such itemsets from F before applying AprioriGen(F)

64

Problem with Closed Set Approach

- Exact Support Equality: Requires supports of some itemsets and their supersets to be **exactly** equal.
- Mushroom Database Example: Addition of 438 tuples to 8,124 tuple database causes number of closed frequent itemsets at $minsup=20\%$ to increase from 1,390 to 15,541 – 11 times!

Talk Outline

- Introduction
- First-time Mining of Association Rules
- Conciseness of Mining Results
- Conclusions

Architecture for Association Rule mining

