

## Data Warehousing

## Roadmap

- What is a data warehouse, data mart
- Multi-dimensional data modeling
- Data warehouse design – schemas, indices
- The Data Cube operator – semantics and computation
- Aggregate View Selection

## What is Data Warehouse?

- Definition: Collection of *decision support* technologies to enable *knowledge worker* (manager, analyst) to make better and faster decisions.
- 4.7 Billion market worldwide [2006 figure, olapreport.com]
  - Retail industries: user profiling, inventory management
  - Financial services: credit card analysis, fraud detection
  - Telecommunications: call analysis, fraud detection
- **Problems:**
  - Takes too long before anything is delivered (6-24 mo.)
  - Costly: hardware, software, manpower, training (>\$1M)
- **Benefits:** Cleaning the muddy wind-shield of a car.

## Data Warehouse Definition

- **Definition (W.H.Inmon):** “A data warehouse is a subject-oriented, integrated, time-variant, and nonvolatile collection of data in support of management's decision-making process.”
- Subject Oriented: Organized around major subjects: *customer*, *product*, *sales*
  - Focus on modeling and analysis of data for decision makers, not on daily operations or transaction processing
- Integrated: Integrate multiple, heterogeneous data sources; exclude data that are not useful in the decision support process
- Time Variant: Needs large time horizon for trend analysis (current and past data)
- Non-Volatile: Physically separate non-volatile store from the operational environment

## Data Marts

---

- A *subset* of what would be in a data warehouse
- Used for a single department, division or geographical location
- Much cheaper than implementing a complete data warehouse
- Can be used as “proof of concept”
- Data marts can co-exist with a data-warehouse
- Multiple data marts should be integrated to ensure consistency and synchronization

## Why not Using Existing DB?

- DBMS is for On Line Transaction Processing (OLTP)
  - automate day-to-day operations (purchasing, banking etc)
- Data Warehouse is for On Line Analytical Processing (OLAP)
  - need historical data for trend analysis


## OLTP vs. OLAP

	OLTP	OLAP
<b>users</b>	clerk, IT professional	knowledge worker
<b>function</b>	day to day operations	decision support
<b>DB design</b>	application-oriented	subject-oriented
<b>data</b>	current, up-to-date detailed, flat relational isolated	historical, summarized, multidimensional integrated, consolidated
<b>usage</b>	repetitive	ad-hoc
<b>access</b>	read/write index/hash on prim. key	lots of scans
<b>unit of work</b>	short, simple transaction	complex query
<b># records accessed</b>	tens	millions
<b>#users</b>	thousands	hundreds
<b>DB size</b>	100MB-GB	100GB-TB
<b>metric</b>	transaction throughput	query throughput, response

7

## Data Warehouse Architecture

- Extract data from operational data sources
  - clean, transform
- Bulk load/refresh
  - warehouse is offline
- OLAP-server provides multidimensional view



## Examples of OLAP

- Comparisons (this period v.s. last period)
  - Show me the sales per store for this year and compare it to that of the previous year to identify discrepancies
- Ranking and statistical profiles (top N/bottom N)
  - Show me sales, profit and average call volume per day for my 10 most profitable salespeople
- Custom consolidation (market segments, ad hoc groups)
  - Show me an abbreviated income statement by quarter for the last four quarters for my northeast region operations


## Roadmap

1. What is the data warehouse, data mart
2. Multi-dimensional data modeling
3. Data warehouse design – schemas, indices
4. The Data Cube operator – semantics and computation
5. Aggregate View Selection

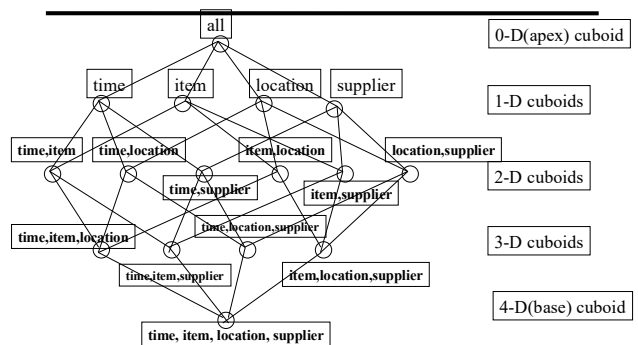
10

## Multidimensional Modeling

- Example: compute total sales volume per *product* and *store*




## Cube: A Lattice of Cuboids



12


## Dimensions and Hierarchies

- A cell in the cube may store values (measurements) relative to the combination of the labeled dimensions



## Common OLAP Operations

- **Roll-up:** move up the hierarchy
  - e.g. given total sales per city, we can roll-up to get sales per state
- **Drill-down:** move down the hierarchy
  - more fine-grained aggregation




## Pivoting

- **Pivoting:** aggregate on selected dimensions
  - usually 2 dims (cross-tabulation)

Sales		Product				
		1	2	3	4	ALL
Store	1	454	-	-	925	1379
	2	468	800	-	-	1268
	3	296	-	240	-	536
	4	652	-	540	745	1937
	ALL	1870	800	780	1670	5120

## Slice and Dice Queries

- **Slice and Dice:** select and project on one or more dimensions



## Roadmap

1. What is the data warehouse, data mart
2. Multi-dimensional data modeling
3. Data warehouse design – schemas, indices
4. The Data Cube operator – semantics and computation
5. Aggregate View Selection

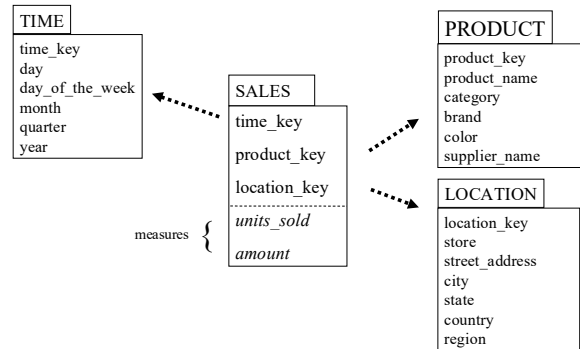
## Data Warehouse Design

- **ROLAP**
  - Store data in RDBMS
  - Provide a multi-dimensional view of this data
  - Makes use of existing technology
  - Products: Redbrick, Informix, Sybase, SQL server
- **MOLAP**
  - Directly implement multi-dimensional model
  - Uses arrays
  - Lots of compression for sparse arrays
  - Products: Essbase, Oracle Express

## ROLAP Schemas

- Most data warehouses adopt a *star schema* to represent the multidimensional model
- Each dimension is represented by a *dimension-table*
  - LOCATION(location\_key, store, street\_address, city, state, country, region)
  - dimension tables are **not normalized**
- Transactions are described through a *fact-table*
  - each tuple consists of a pointer to each of the dimension-tables (foreign-key) and a list of measures (e.g. sales)
- Snowflake Schema:** Same as above, but dimension tables are normalized to provide explicit support for attribute hierarchies

## The Star Schema



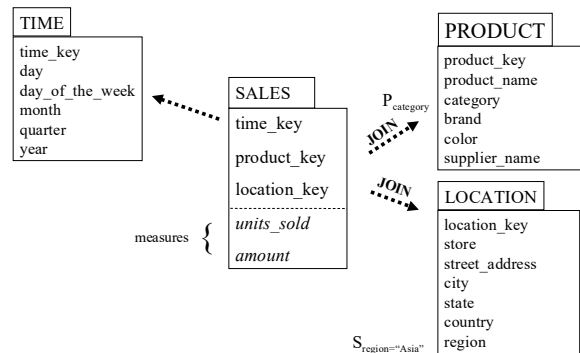
## Advantages of Star Schema

- Facts and dimensions are clearly depicted
  - dimension tables are relatively static, data is loaded (append mostly) into fact table(s)
  - easy to comprehend (and write queries)

"Find total sales per product-category in our stores in Asia"

```
SELECT PRODUCT.category, SUM(SALES.amount)
FROM SALES, PRODUCT, LOCATION
WHERE SALES.product_key = PRODUCT.product_key
AND SALES.location_key = LOCATION.location_key
AND LOCATION.region="Asia"
GROUP BY PRODUCT.category
```

## Star Schema Query Processing




## Indexing OLAP Data: Bitmap Index

- Index on a particular column
- Each value in the column has a bit vector: **bit-op is fast**
- The length of the bit vector:** # of records in the base table
- The *i*-th bit is set if the *i*-th row of the base table has the value for the indexed column
- not suitable for high cardinality domains**

Base table			Index on Region				Index on Type		
Cust	Region	Type	RecID	Asia	Europe	America	RecID	Retail	Dealer
C1	Asia	Retail	1	1	0	0	1	1	0
C2	Europe	Dealer	2	0	1	0	2	0	1
C3	Asia	Dealer	3	1	0	0	3	0	1
C4	America	Retail	4	0	0	1	4	1	0
C5	Europe	Dealer	5	0	1	0	5	0	1

## Join-Index

- Join index relates the values of the dimensions of a star schema to rows in the fact table.
  - a join index on **region** maintains for each distinct region a list of ROW-IDs of the tuples recording the sales in the region
- Join indices can span multiple dimensions OR
  - can be implemented as bitmap-indexes (per dimension)
  - use bit-op for multiple-joins



## Unresolved: Coarse-grain Aggregations

- "Find total sales per product-category in our stores in Asia"
  - Join-index will prune ¾ of the data (uniform sales), but the remaining ¼ is still large (several millions transactions)
    - Index is unclustered
- High-level aggregations are expensive!!!!

LOCATON  
region  
country  
state  
city  
store

⇒ Long Query Response Times

⇒ Pre-computation is necessary

## Unresolved: Multiple Simultaneous Aggregates

4 Group-bys here:

(store,product)  
(store)  
(product)  
( )

Need to write 4 queries!!!

Cross-Tabulation (products/store)

Sales		Product				
		1	2	3	4	ALL
Store	1	454	-	-	925	1379
	2	468	800	-	-	1268
	3	296	-	240	-	536
	4	652	-	540	745	1937
	ALL	1870	800	780	1670	5120

Sub-totals per store (rows 1-4)  
Sub-totals per product (columns 2-5)  
Total sales (bottom-right cell)

## Roadmap

1. What is the data warehouse, data mart
2. Multi-dimensional data modeling
3. Data warehouse design – schemas, indices
4. The Data Cube operator – semantics and computation
5. Aggregate View Selection

27


## The Data Cube Operator (Gray et al)

- All previous aggregates in a single query:

```
SELECT LOCATION.store, SALES.product_key, SUM (amount)
FROM SALES, LOCATION
WHERE SALES.location_key=LOCATION.location_key
CUBE BY SALES.product_key, LOCATION.store
OR
CUBE product_key, store BY SUM(SALES.amount)
```

Challenge: Optimize Aggregate Computation

## Data Cube: Multidimensional View



## Relational View of Data Cube

Sales		Product				
		1	2	3	4	ALL
Store	1	454	-	-	925	1379
	2	468	800	-	-	1268
	3	296	-	240	-	536
	4	652	-	540	745	1937
	ALL	1870	800	780	1670	5120

Store	Product_key	sum(amount)
1	1	454
1	4	925
2	1	800
2	4	468
3	1	296
3	3	240
4	1	625
4	3	240
4	4	745
ALL	ALL	1379
ALL	ALL	1268
ALL	ALL	536
ALL	ALL	1937
ALL	ALL	1870
ALL	ALL	800
ALL	ALL	780
ALL	ALL	1670
ALL	ALL	5120

```
SELECT LOCATION.store, SALES.product_key, SUM (amount)
FROM SALES, LOCATION
WHERE SALES.location_key=LOCATION.location_key
CUBE BY SALES.product_key, LOCATION.store
```


## Other Extensions to SQL

- Complex aggregation at multiple granularities (Ross et. al. 1998)
    - Compute *multiple dependent* aggregates
- ```

SELECT LOCATION.store, SALES.product_key, SUM (amount)
FROM SALES, LOCATION
WHERE SALES.location_key=LOCATION.location_key
CUBE BY SALES.product_key, LOCATION.store: R
SUCH THAT R.amount = max(amount)
  
```
- Other proposals: the MD-join operator (Chatziantoniou et. al. 1999)


## Data Cube Computation

- Model dependencies among the aggregates:




## Computation Directives

- Hash/sort based methods (Agrawal et. al. VLDB'96)
  - Smallest-parent**
  - Cache-results**
  - Amortize-scans**
  - Share-sorts**
  - Share-partitions**



## Alternative Array-based Approach

- Model data as a sparse multidimensional array
  - partition array into chunks (a small sub-cube which fits in memory).
  - fast addressing based on (chunk\_id, offset)
- Compute aggregates in "multi-way" by visiting cube cells in the order which minimizes the # of times to visit each cell, and reduces memory access and storage cost.



**What is the best traversing order to do multi-way aggregation?**

## Roadmap

- What is the data warehouse, data mart
- Multi-dimensional data modeling
- Data warehouse design
  - the star schema, bitmap indexes
- The Data Cube operator
  - semantics and computation
- Aggregate View Selection

## Views and Decision Support



- OLAP queries are typically aggregate queries.
  - Pre-computation is essential** for interactive response times.
  - The CUBE is in fact a collection of aggregate queries, and pre-computation is especially important: lots of work on what is best to pre-compute given a limited amount of space to store pre-computed results.
- Warehouses can be thought of as a collection of asynchronously replicated tables and periodically maintained views.
  - Has renewed interest in view maintenance!

## View Modification (Evaluate On Demand)

**View**

```
CREATE VIEW RegionalSales(category,sales,state)
AS SELECT P.category, S.sales, L.state
FROM Products P, Sales S, Locations L
WHERE P.pid=S.pid AND S.locid=L.locid
```

**Query**

```
SELECT R.category, R.state, SUM(R.sales)
FROM RegionalSales AS R GROUP BY R.category, R.state
```

**Modified Query**

```
SELECT R.category, R.state, SUM(R.sales)
FROM (SELECT P.category, S.sales, L.state
FROM Products P, Sales S, Locations L
WHERE P.pid=S.pid AND S.locid=L.locid) AS R
GROUP BY R.category, R.state
```

37

## View Materialization (Pre-computation)

- Suppose we pre-compute RegionalSales and store it with a clustered B+ tree index on [category,state,sales].
  - Then, previous query can be answered by an index-only scan.

```
SELECT R.state, SUM(R.sales)
FROM RegionalSales R
WHERE R.category="Laptop"
GROUP BY R.state
```

```
SELECT R.state, SUM(R.sales)
FROM RegionalSales R
WHERE R.state="Wisconsin"
GROUP BY R.category
```

Index on pre-computed view is less useful (must scan entire leaf level).

38

## Materialized Views

- A view whose tuples are stored in the database is said to be materialized.
  - Provides fast access, like a (very high-level) cache.
  - Need to maintain the view as the underlying tables change.
  - Ideally, we want incremental view maintenance algorithms.
- Close relationship to data warehousing, OLAP, (asynchronously) maintaining distributed databases, checking integrity constraints, and evaluating rules and triggers.

39


## Issues in View Materialization

- Algorithm to maintain a materialized view?
- What views should we materialize, and what indexes should we build on the pre-computed results?
- Given a query and a set of materialized views (possibly with some indexes), can we use the materialized views to answer the query?

40

## View Selection Problem

- Use some notion of *benefit* per view
- Limit: *disk space*




Hanarayan et al SIGMOD'96:

$$B(v,S) = \sum_{u: u \leq v, C_v(u) < C_S(u)} (C_S(u) - C_v(u))$$

Pick views greedily until space is filled

41

## Reality check: too many views!



- 2^n views for n dimensions (no hierarchies)
- Storage/update-time explosion
- More pre-computation doesn't mean better performance!!!!

[note: original graph at <http://www.olapreport.com>]

42

## Problem Generalization

- Materialize and maintain the right subset of views with respect to the workload and the available resources
- What is the workload?
  - “Farmers” v.s. “Explorers” [Inmon99]
  - Pre-compiled queries (report generating tools, data mining)
  - Ad-hoc analysis (unpredictable)
- What are the resources?
  - Disk space (getting cheaper)
  - Update window (getting smaller)

43

## View Selection Problem

- Selection is based on a workload estimate (e.g. logs) and a given constraint (disk space or update window)
- NP-hard, optimal selection can not be computed > 4-5 dimensions
  - greedy algorithms (e.g. [Harinarayan96]) run at least in polynomial time in the number of views i.e exponential in the number of dimensions!!!
- Optimal selection can not be approximated [Karloff99]
  - greedy view selection can behave arbitrary bad
- Alternatives: use query result caching techniques and reuse prior computations

44

## View Maintenance

- Two steps:
  - Propagate: Compute changes to view when data changes.
  - Refresh: Apply changes to the materialized view table.
- Maintenance policy: Controls when we do refresh.
  - Immediate: As part of the transaction that modifies the underlying data tables. (+ Materialized view is always consistent; - updates are slowed)
  - Deferred: Some time later, in a separate transaction. (- View becomes inconsistent; + can scale to maintain many views without slowing updates)

45

## Deferred Maintenance

- Three flavors:
  - Lazy: Delay refresh until next query on view; then refresh before answering the query.
  - Periodic (Snapshot): Refresh periodically. Queries possibly answered using outdated version of view tuples. Widely used, especially for asynchronous replication in distributed databases, and for warehouse applications.
  - Event-based: E.g., Refresh after a fixed number of updates to underlying data tables.

46

## Sources of Information - Books

1. The Data Warehouse Toolkit – Ralph Kimball ISBN 0-471-15337-0
2. The Data Warehouse Lifecycle Toolkit – Ralph Kimball, Laura Reeves, Warren Thornthwaite & Margy Ross ISBN 0-471-25547-5
3. Data Warehouse Design Solutions – Christopher Adamson & Michael Venerable ISBN 0-471-25195-X

## Sources of Information – Web Sites

Technology Guides for Data Warehousing - [www.techguide.com](http://www.techguide.com)  
Ralph Kimball Associates Articles - [www.ralphkimball.com/html/articles.html](http://www.ralphkimball.com/html/articles.html)  
Data Warehousing - Data Warehousing Knowledge Center - [www.datawarehousing.org](http://www.datawarehousing.org)  
The Data Warehousing Information Center - [www.dwinfocenter.org](http://www.dwinfocenter.org)  
Documenting data replication and data transformation sites on the Net - [www.datawarehousing.com](http://www.datawarehousing.com)  
DM Review Business Intelligence & Data Warehousing Enabling E-Business - [www.dmreview.com/](http://www.dmreview.com/)  
The Data Warehousing Institute - [www.dw-institute.com](http://www.dw-institute.com)  
Intelligent Enterprise Magazine - [www.intelligententerprise.com/](http://www.intelligententerprise.com/)  
DataWarehousing Forum - [www.datawarehousing.com/forum/](http://www.datawarehousing.com/forum/)



## Sources of Information - ListServer

### **Subscribing:**

To subscribe to the data warehousing list server:

- Send an E-mail to [dwlist-request@datawarehousing.com](mailto:dwlist-request@datawarehousing.com)
- The first line of your message must be "subscribe".
- No subject line is required.

### **Unsubscribing:**

To unsubscribe from the data warehousing list server:

- Send an E-mail to [dwlist-request@datawarehousing.com](mailto:dwlist-request@datawarehousing.com)
- The first line of your message must be "unsubscribe".
- No subject line is required.

### **Posting Messages:**

- To post a message, send your message by email to [dwlist@datawarehousing.com](mailto:dwlist@datawarehousing.com)

The End

Thank you!